




Implementing Visual Search with Vocabulary Trees




Mohamed Harmanani

Feature Extraction

BRISK Detector



SIFT Detector



Code to extract SIFT features



```
def get_SIFT_descriptors(img):  
    sift = cv2.SIFT_create()  
    keypoints, descriptors = sift.detectAndCompute(img.copy(), None)  
    kp_list = []  
    for kp in keypoints:  
        kp_list += [(kp.response, kp.pt, kp.size)]  
    return kp_list, descriptors
```

Code to extract BRISK features

```
def get_BRISK_descriptors(img):  
    brisk = cv2.BRISK_create()  
    keypoints, descriptors = brisk.detectAndCompute(img.copy(), None)  
    kp_list = []  
    for i in range(len(keypoints)):  
        kp = keypoints[i]  
        kp_list += [(kp.response, kp.pt, kp.size, i)]  
    return kp_list, descriptors
```

Homography Mapping

Step 1: Keypoint matching



Using SIFT for keypoint matching


```
def match_keypoints(img1, img2, algo='SIFT', thresh=0.45):
    kps_1, desc_1 = _algo_map[algo](img1)
    kps_2, desc_2 = _algo_map[algo](img2)

    best_matches = []
    for i in range(len(desc_1)):
        src, desc1 = kps_1[i][1], desc_1[i]
        matches = []
        for j in range(len(desc_2)):
            dst, desc2 = kps_2[j][1], desc_2[j]
            sim = np.linalg.norm(desc1-desc2, 2)
            matches += [(float(sim), dst, desc2)]

        # get closest match
        indices = [y[0] for y in matches]
        idx = np.argmin(np.array(indices))
        closest = matches[idx]

        # get second closest
        matches_without_closest = matches[:]
        matches_without_closest.remove(closest)
        indices = [y[0] for y in matches_without_closest]
        idx2 = np.argmin(np.array(indices))
        sec_closest = matches_without_closest[idx2]

        # compute phi
        phi = np.linalg.norm(desc1 - closest[2], 2) / np.linalg.norm(desc1 - sec_closest[2], 2)
        if phi <= thresh:
            best_matches += [(phi, (src, closest[1]))]

    return best_matches
```


Step 2: RANSAC

```
def RANSAC_fit_homography(matches, n=5000, thresh=10):
    H = np.zeros((8,9))
    inliers_map = {}
    for i in range(n):
        # Sample 3 matches at random
        sampled = random.sample(matches, 4)
        points_from = [sample[0] for sample in sampled]
        points_to = [sample[1] for sample in sampled]

        # Compute the homography from our samples
        H = get_homography_transform(points_from, points_to)

        # Count and save the number of inliers
        inliers = count_inliers(matches, H, thresh=thresh)
        num_inliers = len(inliers)
        inliers_map[num_inliers] = inliers_map.setdefault(num_inliers, []) + [H]

    idx = max(list(inliers_map.keys()))
    return inliers_map[idx][-1]
```



Python implementation of RANSAC

RANSAC Homography Result



Warped representation for example image


Step 3: Superpose Image



Superposing warped representations onto test images

Efficient Index Retrieval

Building a feature database

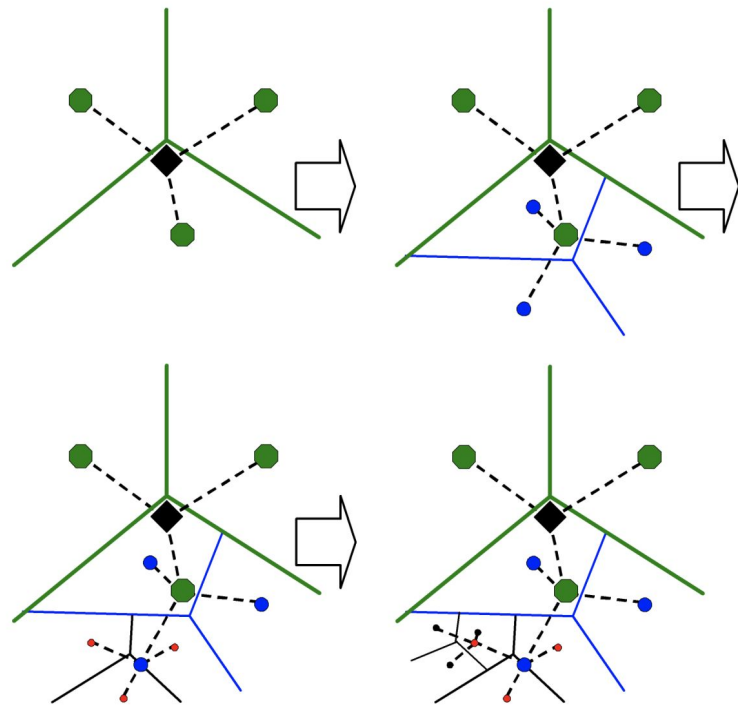


```
[244, 255, 239, ..., 13, 25, 48],  
[244, 255, 255, ..., 12, 8, 253],  
[252, 255, 255, ..., 228, 238, 207],  
...,  
[235, 191, 247, ..., 255, 159, 63],  
[160, 251, 159, ..., 239, 127, 14],  
[132, 123, 223, ..., 239, 175, 143]],
```

Finalized database of features

Partition the dataset using k-means

```
{3: array([[244, 255, 239, ..., 13, 25, 48],
          [244, 255, 255, ..., 12, 8, 253],
          [160, 255, 255, ..., 255, 255, 255],
          ...,
          [180, 25, 251, ..., 223, 159, 63],
          [252, 25, 64, ..., 13, 0, 0],
          [160, 251, 159, ..., 239, 127, 14]], dtype=uint8),
 0: array([[252, 255, 255, ..., 228, 238, 207],
          [252, 191, 255, ..., 225, 230, 207],
          [160, 255, 255, ..., 13, 128, 130],
          ...,
          [176, 123, 207, ..., 64, 226, 231],
          [235, 191, 247, ..., 255, 159, 63],
          [132, 123, 223, ..., 239, 175, 143]], dtype=uint8),
 2: array([[255, 159, 35, ..., 140, 153, 48],
          [160, 255, 255, ..., 0, 129, 2],
          [252, 255, 239, ..., 141, 153, 27],
          ...,
          [180, 120, 223, ..., 239, 159, 59],
          [180, 59, 223, ..., 255, 255, 49],
          [ 52, 64, 8, ..., 62, 31, 16]], dtype=uint8),
 1: array([[249, 191, 231, ..., 8, 17, 50],
          [207, 255, 239, ..., 12, 25, 8],
          [128, 255, 223, ..., 22, 230, 205],
          ...,
          [ 0, 122, 143, ..., 0, 196, 230],
          [ 3, 96, 28, ..., 176, 241, 243],
          [ 75, 254, 239, ..., 0, 0, 0]], dtype=uint8)}
```



```
def build_tree(self, features):
    centroid = self.get_cluster_centroid(features)
    self.set_root_value(centroid)

    # Find centroid, set root, and eliminate corresponding row
    for i in range(features.shape[0]):
        if np.all(features[i] == centroid):
            features = np.delete(features, i, axis=0)
            break

    if self.L <= 0:
        return

    if features.shape[0] < self.bf:
        for c in range(features.shape[0]):
            child_id = INDICES.pop(INDICES.index(min(INDICES)))  #(self.K * self.id) + (c + 1)
            self.children[c] = VocabTree(self.bf, self.L - 1, self.k_means_func, features, child_id)
            self.children[c].set_root_value(features[c])
        return

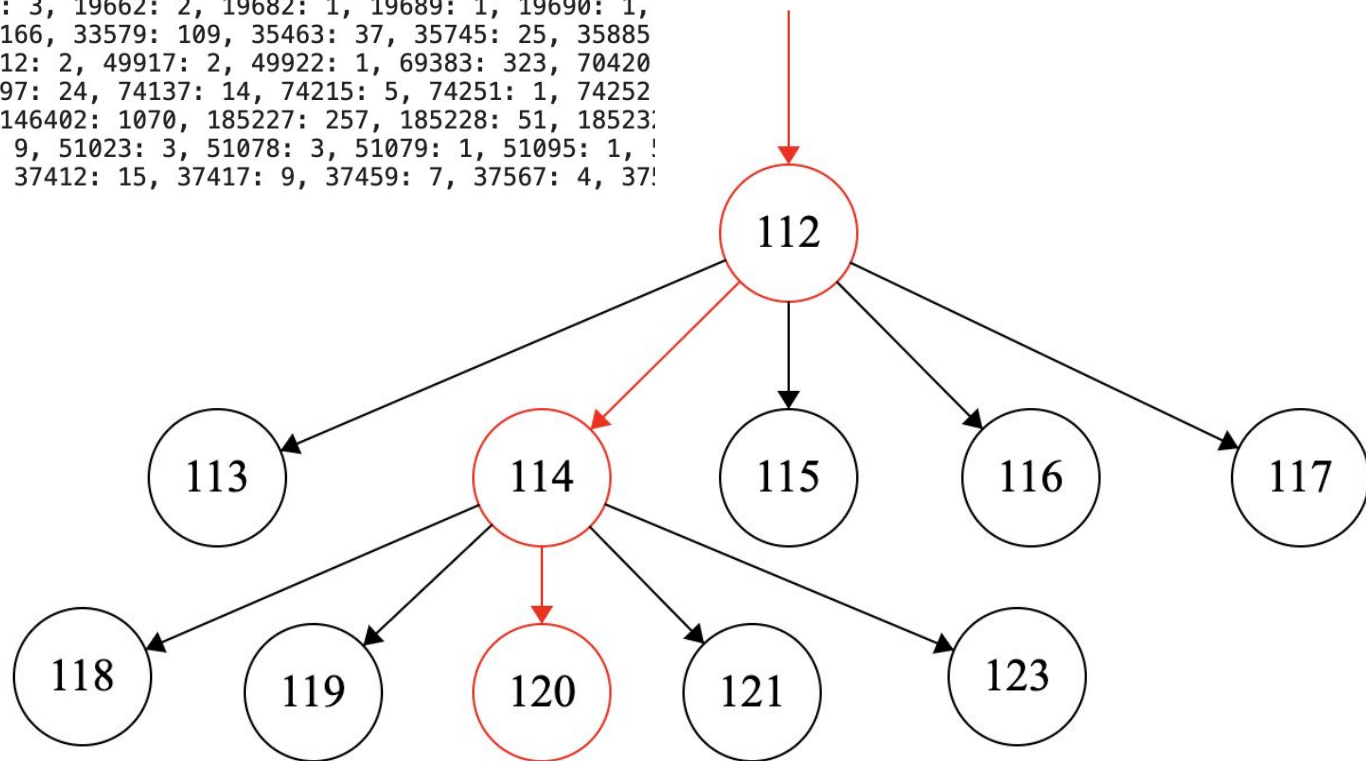
    labels = self.KMeans.fit_predict(features)
    clusters = self.partition_dataset(features, labels)

    for c in range(len(self.children)):
        if c in labels: # Not always possible to have as many clusters as we want
            child_id = INDICES.pop(INDICES.index(min(INDICES)))
            self.children[c] = VocabTree(self.bf, self.L - 1, self.k_means_func, clusters[c], child_id)
            self.children[c].build_tree(clusters[c])
```


Propagate features down the vocabulary tree

```
'{37968: 559, 44645: 63, 44646: 23, 46483: 11, 46823: 9, 46824: 4, 46851: 3, ,  
238, 19202: 101, 19660: 12, 19661: 3, 19662: 2, 19682: 1, 19689: 1, 19690: 1,  
301: 1, 2309: 1, 2314: 1, 30732: 166, 33579: 109, 35463: 37, 35745: 25, 35885  
49349: 7, 49758: 3, 49911: 2, 49912: 2, 49917: 2, 49922: 1, 69383: 323, 70420  
73108: 3, 73109: 2, 73110: 1, 73397: 24, 74137: 14, 74215: 5, 74251: 1, 74252  
1, 24916: 4, 24945: 2, 24947: 2, 146402: 1070, 185227: 257, 185228: 51, 18523:  
3: 1, 185435: 1, 50198: 9, 50687: 9, 51023: 3, 51078: 3, 51079: 1, 51095: 1, !  
1, 48977: 1, 48984: 1, 36475: 25, 37412: 15, 37417: 9, 37459: 7, 37567: 4, 37!  
1, 146403: 396, '
```

Bag of features



Code


```
def _propagate_feature(self, f, code):
    dist_map = {}
    id_map = {}
    for i in range(len(self.children)):
        c = self.children[i]
        if not c:
            return
        dist = np.linalg.norm(f - c.value)
        dist_map[dist] = c.id
        id_map[c.id] = i
    min_id = dist_map[min(dist_map.keys())]
    code[min_id] = code.setdefault(min_id, 0) + 1
    if not self.children[id_map[min_id]].is_leaf_node():
        self.children[id_map[min_id]]._propagate_feature(f, code)
```

```
def compute_path(self, image):
    code = {}
    _, desc = get_SIFT_descriptors(image.copy())
    for i in range(desc.shape[0]):
        #path = []
        self._propagate_feature(desc[i], code)
        #for j in range(len(path)):
        #    enc = (path[j] + i) * j
        #    code[enc] = code.setdefault(enc, 0) + 1
    return code
```

tf-idf weighting

```
{37968: 559, 44645: 63, 44646: 23, 46483: 11, 46823: 9, 46824: 4, 46851: 3, 46853: 2, 46861: 1, 0: 917, 11824: 547, 16950: 238, 19202: 101, 19660: 12, 19661: 3, 19662: 2, 19682: 1, 19689: 1, 19690: 1, 1: 179, 2: 53, 1783: 19, 2264: 11, 2265: 1, 2301: 1, 2309: 1, 2314: 1, 30732: 166, 33579: 109, 35463: 37, 35745: 25, 35885: 7, 35895: 5, 35896: 3, 35899: 3, 47240: 23, 49349: 7, 49758: 3, 49911: 2, 49912: 2, 49917: 2, 49922: 1, 69383: 323, 70420: 87, 72803: 23, 73082: 6, 73102: 4, 73103: 3, 73108: 3, 73109: 2, 73110: 1, 73397: 24, 74137: 14, 74215: 5, 74251: 1, 74252: 1, 74253: 1, 22673: 80, 24676: 60, 24677: 1, 24916: 4, 24945: 2, 24947: 2, 146402: 1070, 185227: 257, 185228: 51, 185232: 21, 185402: 14, 185403: 2, 185429: 1, 185433: 1, 185435: 1, 50198: 9, 50687: 9, 51023: 3, 51078: 3, 51079: 1, 51095: 1, 51097: 1, 48290: 4, 48820: 2, 48831: 2, 48941: 1, 48977: 1, 48984: 1, 36475: 25, 37412: 15, 37417: 9, 37459: 7, 37567: 4, 37587: 4, 37600: 3, 37601: 3, 37605: 1, 73113: 1, 146403: 396, '}
```

Bag of features



```
array( [-5083.73743971, -726.20099085, -154.50654109, ...,  
        0., 0., 0.] )
```

Finalized tf-idf representation

```
def compute_weight_vector(self, db_imgcodes):
    w = np.zeros((self.num_nodes, ), dtype=np.float64)
    imgmatrix = self.build_image_matrix(db_imgcodes)
    N = imgmatrix.shape[0]
    Ni = 0
    for i in range(imgmatrix.shape[1]):
        codevec = imgmatrix[:, i]
        Ni = np.sum(codevec)

        if Ni == 0:
            w[i] = 0
        else:
            w[i] = np.log(N/Ni)

    for imgname in self.imgdb:
        self.imgdb[imgname] = w * self._vec_from_dict(self.imgdb[imgname])



    return w
```

Computing the weights



```
def compute_img_code(self, image):
    imgcode = self.compute_path(image)
    db_imgcodes = self.compute_img_paths()
    w = self.compute_weight_vector(db_imgcodes)
    return w * self._vec_from_dict(imgcode)
```

Applying the weights

Retrieving matches



```
[ './DVDcovers/the_incredibles.jpg',  
  './DVDcovers/the_factory.jpg',  
  './DVDcovers/argo.jpg',  
  './DVDcovers/non_stop.jpg',  
  './DVDcovers/reservoir_dogs.jpg',  
  './DVDcovers/10_things_i_hate_about_y',  
  './DVDcovers/the_santa_clause.jpg',  
  './DVDcovers/the_last_stand.jpg',  
  './DVDcovers/matrix.jpg',  
  './DVDcovers/the_good_wife_season_3.j
```



```
[ './DVDcovers/always.jpg']  
[ './DVDcovers/matrix.jpg']  
[ './DVDcovers/outlander.jpg']  
[ './DVDcovers/man_of_steel.jpg']  
[ './DVDcovers/the_sum_of_all_fears.jp  
[ './DVDcovers/reservoir_dogs.jpg']  
[ './DVDcovers/legends_of_the_fall.jpg  
[ './DVDcovers/the_grey.jpg']  
[ './DVDcovers/sommersby.jpg']  
[ './DVDcovers/supernatural.jpg']
```

Code

```
imgmap = {}  
for img in os.listdir('./DVDcovers/'):   
    imgname = './DVDcovers/' + img  
    try:  
        im = io.imread(imgname)  
    except:  
        continue  
    x = (vt.imgdb[imgname])  
    r = cos_sim(x, code_mx)  
    imgmap[r] = imgmap.setdefault(r, []) + [imgname]  
  
for u in sorted(imgmap)[::-1][:10]:  
    print(imgmap[u])
```

Code for printing top 10 retrieved matches


Issues & Discussion

$$O(K^L)$$

Problem Images




Score: ~0.79



Score: ~0.94

Problem Images



Excessive Noise

